



designed by  freepik.com

# XSI (system V) IPC

# System V IPC

- UNIX specification has undergone many stages
- Today, the Open Group extends some more UNIX definitions, and these are in the SUSv4 (Single Unix Specifications version 4) which is the current “UNIX” standard.
- Still, the following IPC objects were first implemented in System V unix, and are still called that way

# identifiers and keys

- Each IPC structure (message queue, semaphore, or shared memory segment) in the kernel is referred to by a non-negative integer identifier
- Unlike file descriptors, IPC identifiers are not small integers. Their numbers are not re-used if deleted
- Keys are used by processes to rendezvous with IPC objects. Users can use meaningful names with keys.

# keys basics

- Whenever we create a System V object (msgget, shmget, semget) we have to specify the key.
- The key is of type `key_t` (generally defined as a long int)
- The key is used to create the identifier.
- If we set key to `IPC_PRIVATE`, then the creator can create a new IPC object, and can share it with other processes related by `fork()` (children, grandchildren..) or by storing the identifier in a common file.

# using keys(1) : IPC\_PRIVATE

- If we set key to IPC\_PRIVATE, then the caller will create a new IPC object
- The caller can store the identifier in a place where some other process can find it (like a file)
- If we have a parent-child situation (via fork()), then the child has an access to the identifier after the fork.

# using keys(2) : agree on a value

- Both sides can agree on a value
- This can be achieved (for example) by using a common header file
- The problem is that the object can already exist
- In this case, the xget function fails (with error EEXIST)
- The creator can try to delete it, and re-create it again

# using keys(3) : using ftok

- Both sides must agree on a **path** and **project id**
- path must be a real accessible file
- Id is an int, but only the lower 8 bits are used
- Note:  
ftok can still create keys that were already used in some situations.  
Always be ready to deal with errors.

## flags: IPC\_CREATE & IPC\_EXCL

- All 3 get functions (msgget, shmget, semget) have a flags field, controlling what is done.
- If we use a “new” key, then we must specify IPC\_CREATE to create a new object.
- If we refer to a used key, we should get an access to an already created object
- If we specify both IPC\_CREATE and IPC\_EXCL
- If we use IPC\_PRIVATE as a key, we will never get access to an old object, we’ll always create a new one.



# flags: permissions

- flags field (badly documented) is also used for permissions:

S_IRUSR	-	read permission, owner
S_IWUSR	-	write permission, owner

S_IRGRP	-	read permission, group
S_IWGRP	-	write permission, group

S_IROTH	-	read permission, others
S_IWOTH	-	write permission, others



designed by  freepik.com

# message queue

# System V Message queue

- A message queue is a linked list of messages stored within the kernel
- identified by a message queue identifier
- New messages are added to the end of a queue
  - by `msgsnd`
- Every message has:
  - a positive long integer type field
  - a non-negative length
  - the actual data bytes

# msqid\_ds

- Each message queue has a data structure containing data about the current status of the queue.
- We can get a copy of it using `msgctl`

```
struct msqid_ds {  
    struct ipc_perm  msg_perm;      /* see Section 15.6.2 */  
    msgqnum_t        msg_qnum;      /* # of messages on queue */  
    msglen_t         msg_qbytes;    /* max # of bytes on queue */  
    pid_t            msg_lspid;     /* pid of last msgsnd() */  
    pid_t            msg_lrpid;     /* pid of last msgrcv() */  
    time_t           msg_stime;     /* last-msgsnd() time */  
    time_t           msg_rtime;     /* last-msgrcv() time */  
    time_t           msg_ctime;     /* last-change time */  
    :  
};
```

# msgget

- **msgget** is used to create a new message queue
- New messages are added to the end of a queue by **msgsnd**
- Every message has a positive long integer **type** field
- Messages are fetched from a queue by **msgrcv**

```
#include <sys/msg.h>
```

```
int msgget(key_t key, int flag);
```

Returns: message queue ID if OK, -1 on error

# msgctl

```
#include <sys/msg.h>
```

```
int msgctl(int msqid, int cmd, struct msqid_ds *buf);
```

Returns: 0 if OK, -1 on error

## ● Set cmd to:

- IPC\_STAT get a copy of msqid\_ds
- IPC\_SET set fields like msg\_perm.uid or mode
- IPC\_RMID Remove the message queue

## msgsnd(cont.)

- msgsnd should block if the queue is full (no more empty slots, or no more available bytes, both are system limits)
- If we specify IPC\_NOWAIT, then instead of blocking, the call returns immediately with error EAGAIN (that means: try AGAIN later)
- If we are blocked, then signal will interrupt the blocking (returning EINTR)
- If we are blocked and the queue is removed, then the block is interrupted (returning EIDRM)

# receive order

- We don't have to fetch the messages in a first-in, first-out order
- Instead, we can fetch messages based on their **type** field



# msgrcv(1)

- Messages are retrieved from a queue by msgrcv.

```
#include <sys/msg.h>
```

```
ssize_t msgrcv(int msqid, void *ptr, size_t nbytes, long type, int flag);
```

Returns: size of data portion of message if OK, -1 on error

- *ptr* should point to the same structure as msgsnd !!! (return message will be stored there)
- *nbytes* has the length of the data we are willing to receive (only the buffer)

# msgrcv(2)

- If message is too large, then:
  - setting MSG\_NOERROR will truncate the message (but we do not know that it was truncated)
  - if not set, we receive an error (E2BIG), and the message stays in the queue

# msgrcv(3)

● The type argument lets us specify which message we want:

- `type == 0` The first message on the queue
- `type > 0` specify exactly the type we look for
- `type < 0` specify a number,  
reading types smaller than this number  
(starting with the lowest)

# msgrcv(4)

- We can specify a flag value of IPC\_NOWAIT
- ENOMSG will say that message (such as we specified) could not be found
- The same behaviour as that of msgsnd regarding removal of the queue or signal



designed by  freepik.com

# MQ exercises

# CW: ping pong - step 1 (learn the api)

- Create 2 programs: Ping and Pong.
- Ping is run first (with -c)
- ping sends a series of messages to pong
- Pong responds to each one  
(How many queues do we need ?)
- Pong does not know how many messages will be sent !!!
- Both sides must exit gracefully, but the queue may stay in place

# CW: ping pong - step 1 (cont.)

## ● Ping parameters:

- c create the queue
- v verbose
- e implement EOF as a type (no value)
- d delete queue
- f queue name (optional)
- n number of messages
- s sleep time (msec)

# CW: ping pong - step 1 (cont.)

## ● Pong parameters:

- v     verbose
- e     implement EOF as a type (no value)
- f     queue name (optional)
- s     sleep time (msec)



## HW: ping pong - step 2

- In step 2 we may run several Ping programs, and several Pong programs, each in a terminal.
- Each Ping may receive a different number of messages (-n parameter).
- At least one Ping should run first.
- ONE QUEUE !!!
- Gracefull exit !!!
- Same parameters



designed by  freepik.com

# System-V semaphores

# semaphores

- A semaphore is a counter used to provide access to a shared data object for multiple processes.
- It could be used to control access to a shared resource
- This is based on an “atomic” operation, containing 2 functions:
  - first test if a condition is set (e.g: `semaphore > 0`)
  - then set the condition (`semaphore --`)
- No other process (or thread) can go between these operations creating a race.

# semaphores roles

● The functionality of a semaphore can be thought of serving 3 separate roles:

- notification: go to sleep until semaphore is set
- counting: “safe” counting
- lock: make sure only one process has access to a resource

# system V Semaphores

- Are created in sets (more than a single semaphore at once)
- Are created (`semget`), then initialized (`semctl`)  
This can be a problem
- We have to remove semaphores after using them

# setget

- creates a set of semaphores:

```
#include <sys/sem.h>
```

```
int semget(key_t key, int nsems, int flag);
```

Returns: semaphore ID if OK, -1 on error

- If we specify a new set, we have to set nsems.  
If referencing an existing set, we can leave it 0.
- The identifier retrieved is an identifier for a set of semaphores (set of values for the semaphore)

# setctl

- Use semctl for many operations for the semaphore set, including (among other):
  - IPC\_RMID , remove the set
  - GETVAL get the value for a single semaphore in the set
  - SETVAL set a value
  - GETNCNT, # of proccesses waiting that  $\text{semval} \geq 0$
  - GETZCNT # of proccesses waiting that  $\text{semval} = 0$

# semop

- Used to do the atomic operations on a set of values.

```
#include <sys/sem.h>
```

```
int semop(int semid, struct sembuf semoparray[], size_t nops);
```

Returns: 0 if OK, -1 on error

one look like this:

```
struct sembuf {  
    unsigned short sem_num; /* member # in set (0, 1, ..., nsems-1) */  
    short          sem_op;  /* operation (negative, 0, or positive) */  
    short          sem_flg; /* IPC_NOWAIT, SEM_UNDO */  
};
```



# sem\_op field

- $\text{sem\_op} > 0$ , adding this to the value of the semaphore (returning resources) not blocking
- $\text{sem\_op} < 0$ , we want this number of units from the semaphore.
  - if  $\text{semval} - \text{sem\_op} \geq 0$  then we do not block, and subtract this number
  - if  $\text{semval} - \text{sem\_op} < 0$ , then we
    - \* block if `IPC_NOWAIT` is not set
    - \* return with error if it is

# waiting to zero

- `sem_op == 0` means that we want to wait (i.e. block) until `semval` will come to zero.
- This is unique, and can be useful sometimes.

# semop is atomic

- semop is atomic !!!
- That means it will do its operations on all values, or on none of them !!!