

System Calls

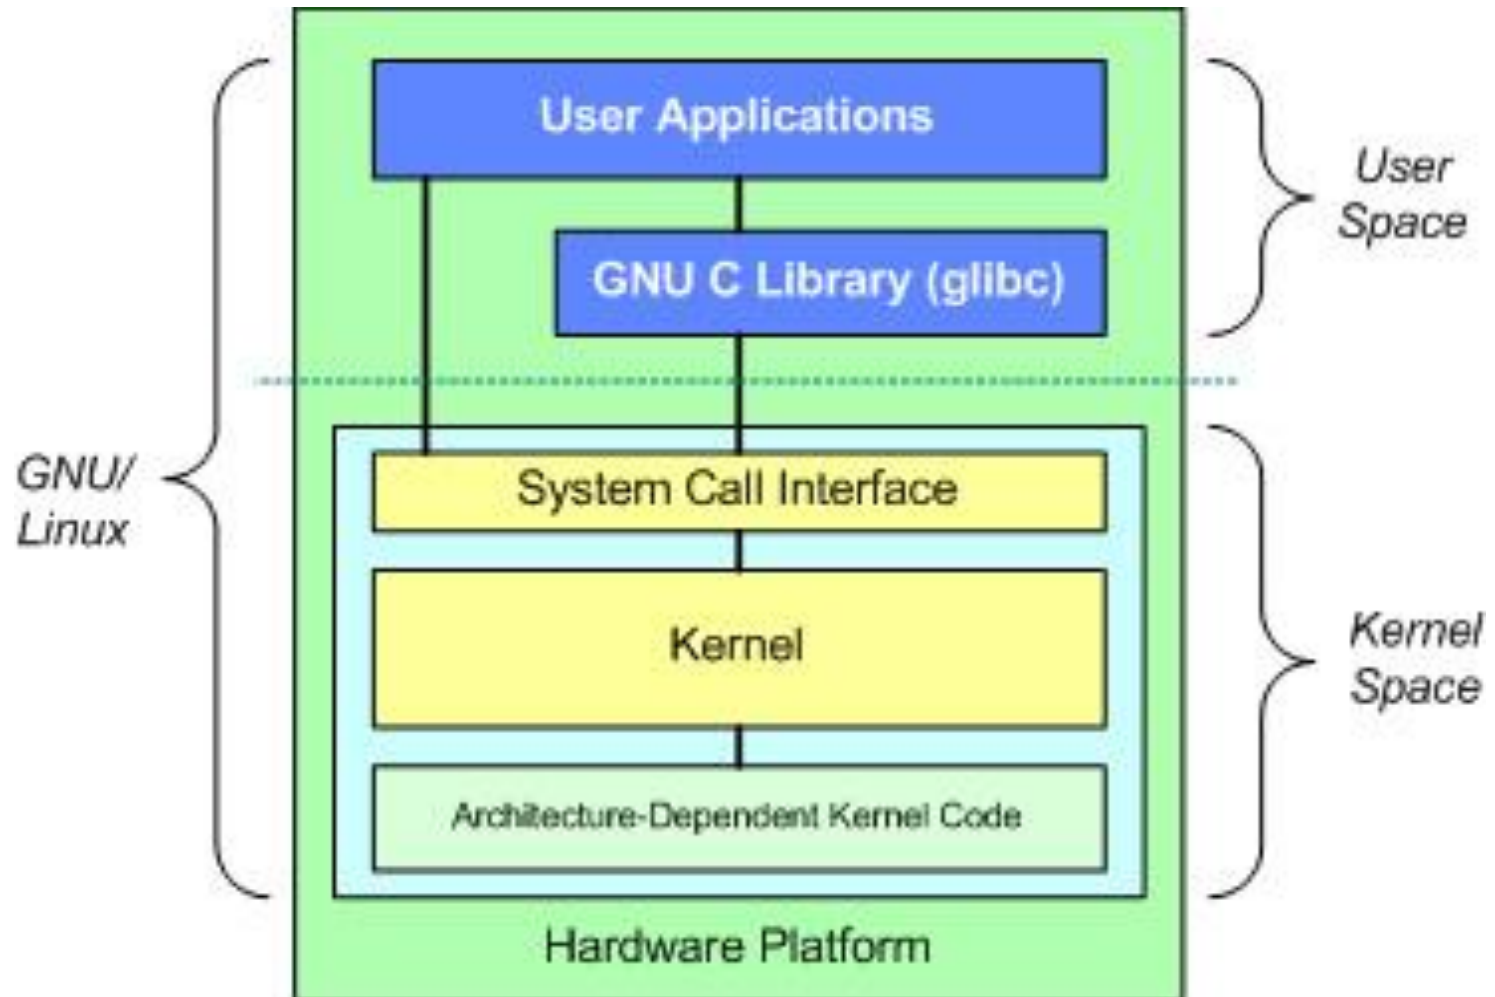
Using Tanenbaum's Modern Operating Systems (3rd edition)

© 2008 Prentice-Hall, Inc. All rights reserved. 0-13-6006639

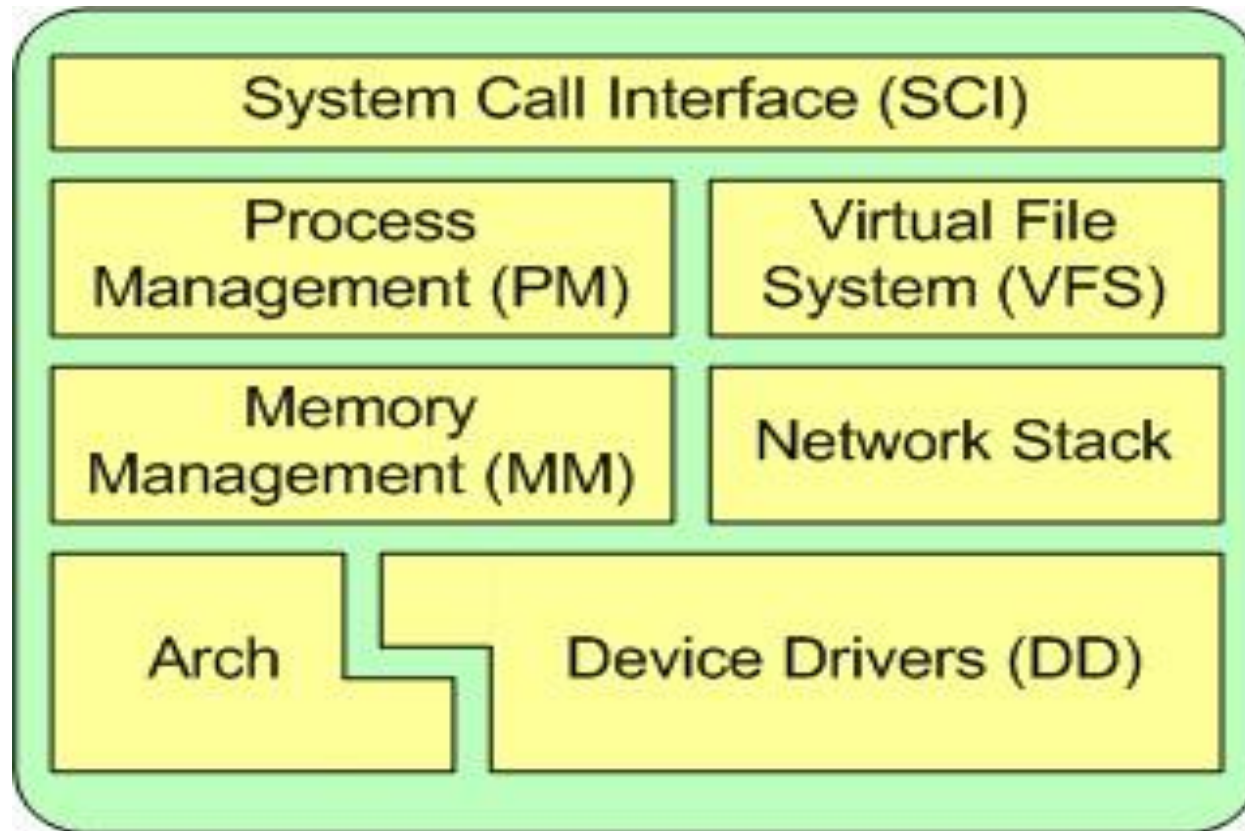
System Call

- What is a system call?
 - API to the OS
 - Requests a service from an operating system's kernel
- Two CPU modes :
 - User mode:
 - Limits the address space of the program.
 - Prevents the application from directly using devices.
 - Kernel mode:
 - All is allowed.
- POSIX defines about 100 system calls:
 - open, read, write, close, wait, execve, fork, exit, kill

Linux OS Architecture



Linux OS Architecture (II)



glibc

- The C standard library provides macros, type definitions, and functions:
 - Basic manipulation:
 - String handling.
 - Mathematical computations.
 - System calls:
 - I/O processing.
 - Memory allocation.
 - Other OS services.

glibc (II)

- By calling a system-call:
 - Sets the parameters and TRAP instruction, which:
 - Switches into kernel mode – transfers control to OS
 - Jumps to a single fixed location
- "*fork*" and "*execve*" are *glibc* functions that in turn call the "fork" and "execve" system-calls.

System Calls (I)

- Process control:
 - load
 - execute
 - create process
 - terminate process
 - get/set process attributes
 - wait for time, wait event, signal event
 - allocate, free memory

System Calls (II)

- File management:
 - create file, delete file
 - open, close
 - read, write, reposition
 - get/set file attributes
- Device management:
 - attach or detach devices
 - read, write, reposition
 - get/set device attributes

System Calls (III)

- Information:
 - get/set time or date
 - get/set system data
 - get/set process, file, or device attributes
- Communications:
 - create, delete communications connection
 - send, receive messages
 - transfer status information
 - attach or detach remote devices

System Call Flow

- C program:


...

```
count = read(fd,buffer,nbytes);
```

...

- counts is the number of read bytes.
 - if failed count=-1 and *errno* is set.
- The library function *read* is called
- The system call number is put into a register
- A TRAP instruction is executed to switch to kernel mode
- The kernel starts at a fixed address
- The kernel executes the system call handler that is a function pointer assigned to the specific system call number
- The handler is executed and control is returned to the user-space library at the instruction after the TRAP

System Call (11 steps)



Blocking System Call

Example:

- The user calls *read()* from the keyboard but nothing has been typed yet.
- What happens?
 - The system call will run the kernel code for blocking the process until a keyboard interrupt arrives
 - In the meantime, another process is taken from the Ready Queue to be executed

POSIX System Calls (I)

Process management

| Call | Description |
|--|--|
| <code>pid = fork()</code> | Create a child process identical to the parent |
| <code>pid = waitpid(pid, &statloc, options)</code> | Wait for a child to terminate |
| <code>s = execve(name, argv, environp)</code> | Replace a process' core image |
| <code>exit(status)</code> | Terminate process execution and return status |

File management

| Call | Description |
|---|---|
| <code>fd = open(file, how, ...)</code> | Open a file for reading, writing, or both |
| <code>s = close(fd)</code> | Close an open file |
| <code>n = read(fd, buffer, nbytes)</code> | Read data from a file into a buffer |
| <code>n = write(fd, buffer, nbytes)</code> | Write data from a buffer into a file |
| <code>position = lseek(fd, offset, whence)</code> | Move the file pointer |
| <code>s = stat(name, &buf)</code> | Get a file's status information |

POSIX System Calls (II)

Directory and file system management

| Call | Description |
|---|--|
| <code>s = mkdir(name, mode)</code> | Create a new directory |
| <code>s = rmdir(name)</code> | Remove an empty directory |
| <code>s = link(name1, name2)</code> | Create a new entry, name2, pointing to name1 |
| <code>s = unlink(name)</code> | Remove a directory entry |
| <code>s = mount(special, name, flag)</code> | Mount a file system |
| <code>s = umount(special)</code> | Unmount a file system |

Miscellaneous

| Call | Description |
|---|---|
| <code>s = chdir(dirname)</code> | Change the working directory |
| <code>s = chmod(name, mode)</code> | Change a file's protection bits |
| <code>s = kill(pid, signal)</code> | Send a signal to a process |
| <code>seconds = time(&seconds)</code> | Get the elapsed time since Jan. 1, 1970 |

Homework + Interview Questions

- What is “system call”?
- Explain in detail how system call is executed
- What is TRAP?